

```

In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 %matplotlib inline
        5 import seaborn as sns
        6 from sklearn.utils.class_weight import compute_class_weight
        7 from sklearn.preprocessing import StandardScaler
        8 from sklearn.linear_model import LogisticRegression
        9 from sklearn.tree import DecisionTreeClassifier
       10 from sklearn.ensemble import RandomForestClassifier
       11 from sklearn.model_selection import train_test_split, GridSearchCV, cr
       12 from sklearn.metrics import accuracy_score, recall_score, precision_sc
       13 from sklearn.metrics import ConfusionMatrixDisplay
       14 from sklearn.metrics import classification_report
       15 from sklearn.pipeline import Pipeline
       16 from imblearn.pipeline import Pipeline as ImbPipeline
       17 from sklearn.decomposition import PCA
       18 from imblearn.over_sampling import SMOTE, BorderlineSMOTE
       19 from google.colab import files
       20 uploaded = files.upload()

```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session.
Please rerun this cell to enable.

Saving pivoted_df.csv to pivoted_df.csv

Load the dataframe in, inspect the data.

```

In [3]: 1 pivoted_df = pd.read_csv('pivoted_df.csv', index_col=0)

```

```

In [4]: 1 pivoted_df.head()

```

Out[4]:

| | season | Age | Throws | Surgery | AB_release_speed_weighted_avg | CH_release_speed_weig |
|---|--------|------|--------|---------|-------------------------------|-----------------------|
| 0 | 2008 | 37.0 | 1 | 0.0 | 0.0 | ε |
| 1 | 2009 | 38.0 | 1 | 0.0 | 0.0 | ε |
| 2 | 2010 | 39.0 | 1 | 0.0 | 0.0 | ε |
| 3 | 2011 | 40.0 | 1 | 0.0 | 0.0 | ε |
| 4 | 2012 | 41.0 | 1 | 0.0 | 0.0 | ε |

5 rows × 130 columns

```

In [5]: 1 pivoted_df.shape

```

Out[5]: (3688, 130)

In [6]: 1 pivoted_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3688 entries, 0 to 3687
Columns: 130 entries, season to SV_vz0_weighted_avg
dtypes: float64(128), int64(2)
memory usage: 3.7 MB
```

In [7]: 1 pivoted_df['Surgery'].value_counts()

```
Out[7]: 0.0    2772
        1.0     916
        Name: Surgery, dtype: int64
```

Time to start modeling! Split target and features and make a baseline model.

In [8]: 1 y = pivoted_df['Surgery']
2 X = pivoted_df.drop('Surgery', axis=1)

In [9]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.

```

In [10]: 1 #Make a pipeline to simplify process
2 logreg_pipeline = Pipeline([
3     ('scale', StandardScaler()),
4     ('logreg', LogisticRegression(solver='liblinear'))
5 ])
6
7 # Define parameter grid to search
8 param_grid = {
9     'logreg__C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
10    'logreg__penalty': ['l1', 'l2'] # Norm used in the penalization
11 }
12
13 # Initialize GridSearchCV with pipeline, parameter grid, and scoring method
14 grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='roc_auc')
15
16 # Assuming X_train and y_train are already defined
17 grid_search.fit(X_train, y_train)
18
19 # Best parameters found
20 print("Best parameters: ", grid_search.best_params_)
21
22 # Best cross-validation score
23 print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
24
25 # Test set score using the best parameters
26 print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))

```

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

Best parameters: {'logreg__C': 10, 'logreg__penalty': 'l1'}

Best cross-validation score: 0.76

Test set score: 0.77

```
In [11]: 1 logreg_pipeline = Pipeline([
2         ('scale', StandardScaler()),
3         ('logreg', LogisticRegression(penalty='l1', C=10.0, solver='liblinear'))
4     ])
```

```
In [12]: 1 logreg_pipeline.fit(X_train, y_train)
```

```
Out[12]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(C=10.0, penalty='l1', solver='liblinear'))])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [13]: 1 logreg_pipeline.score(X_test, y_test)
```

```
Out[13]: 0.7687443541102078
```

```
In [14]: 1 y_pred = logreg_pipeline.predict(X_test)
```

```
In [15]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Baseline LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.79 | 0.93 | 0.86 | 827 |
| 1.0 | 0.59 | 0.28 | 0.38 | 280 |
| accuracy | | | 0.77 | 1107 |
| macro avg | 0.69 | 0.61 | 0.62 | 1107 |
| weighted avg | 0.74 | 0.77 | 0.74 | 1107 |

Dataset is imbalanced, need to adjust. Should also focus on Recall score since this is a medical issue (better to have False Positive than True Negative!)

```
In [16]: 1 y = pivoted_df['Surgery']
2 X = pivoted_df.drop('Surgery', axis=1)
```

```
In [17]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```

In [18]: 1 # Set up pipeline
2 weight_logreg_pipeline = Pipeline([
3     ('scale', StandardScaler()),
4     ('logreg', LogisticRegression(solver='liblinear'))
5 ])
6
7 # Define the parameter grid to search over, including class weights
8 param_grid = {
9     'logreg__C': [0.01, 0.1, 1, 10],
10    'logreg__penalty': ['l1', 'l2'],
11    'logreg__class_weight': [None, 'balanced', {0: 1, 1: 2}, {0: 1, 1: 5}],
12    'logreg__max_iter': [5000],
13    'logreg__tol': [0.01]
14 }
15
16 # Create a scoring function that focuses on recall for the positive class
17 recall_scorer = make_scorer(recall_score, pos_label=1)
18
19 # Initialize GridSearch with pipeline, param grid, and recall
20 grid_search = GridSearchCV(weight_logreg_pipeline, param_grid, cv=5, scoring=recall_scorer)
21
22 # Fit the grid search to the data
23 grid_search.fit(X_train, y_train)
24
25 # Print the best parameters found and the best recall score
26 print("Best parameters: ", grid_search.best_params_)
27 print("Best cross-validation recall score: {:.2f}".format(grid_search.best_score_))
28
29 # Evaluate the best model on the test set
30 best_model = grid_search.best_estimator_
31 y_pred = best_model.predict(X_test)
32 print("Test set recall score: {:.2f}".format(recall_score(y_test, y_pred)))

```

```

Best parameters: {'logreg__C': 0.01, 'logreg__class_weight': {0: 1, 1: 5}, 'logreg__max_iter': 5000, 'logreg__penalty': 'l1', 'logreg__tol': 0.01}
Best cross-validation recall score: 0.87
Test set recall score: 0.85

```

```

In [19]: 1 best_model.fit(X_train, y_train)

```

```

Out[19]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(C=0.01, class_weight={0: 1, 1: 5},
                                                max_iter=5000, penalty='l1',
                                                solver='liblinear', tol=0.01))])

```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [20]: 1 best_model.score(X_test, y_test)

Out[20]: 0.5158084914182475

In [21]: 1 y_pred = best_model.predict(X_test)

In [22]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Weighted LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))




| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.89 | 0.40 | 0.55 | 827 |
| 1.0 | 0.33 | 0.85 | 0.47 | 280 |
| accuracy | | | 0.52 | 1107 |
| macro avg | 0.61 | 0.63 | 0.51 | 1107 |
| weighted avg | 0.75 | 0.52 | 0.53 | 1107 |

Much better model. False Negatives is low, other classes much higher.

```
In [69]: 1 coef = best_model['logreg'].coef_
```

```
In [70]: 1 features = pivoted_df.columns
2
3 zipped = zip(features, coef[0])
4 sorted_pairs = sorted(zipped, key=lambda x: x[1], reverse=True)
5 sorted_pairs
6
7 feature_importances = pd.DataFrame(sorted_pairs, columns=['feature', 'importance'])
8 feature_importances = feature_importances[abs(feature_importances['importance']) > 0.05]
```

```
In [71]: 1 sns.barplot(x='importance', y='feature', data=feature_importances, hue='importance')
2 plt.title('Feature Importances')
3 plt.show()
```



Need to update feature names so they can be understood more easily.

In [72]: 1 feature_importances

Out[72]:

| | feature | importance |
|-----|-------------------------------|------------|
| 0 | Age | 0.350397 |
| 1 | season | 0.204704 |
| 2 | IN_release_pos_x_weighted_avg | 0.106614 |
| 3 | FF_release_pos_y_weighted_avg | 0.056850 |
| 4 | CU_release_pos_x_weighted_avg | 0.028929 |
| 5 | FF_release_pos_z_weighted_avg | 0.015752 |
| 6 | CS_release_speed_weighted_avg | 0.007832 |
| 7 | FF_release_speed_weighted_avg | 0.006973 |
| 8 | ST_vy0_weighted_avg | 0.005001 |
| 9 | CH_vz0_weighted_avg | 0.004601 |
| 117 | FF_vy0_weighted_avg | -0.012434 |
| 118 | ST_release_speed_weighted_avg | -0.013466 |
| 119 | Throws | -0.018221 |
| 120 | CS_vy0_weighted_avg | -0.018479 |
| 121 | PO_release_pos_x_weighted_avg | -0.020661 |
| 122 | ST_release_pos_z_weighted_avg | -0.024097 |
| 123 | ST_release_pos_y_weighted_avg | -0.029266 |
| 124 | KN_release_pos_y_weighted_avg | -0.033056 |
| 125 | FC_release_pos_y_weighted_avg | -0.036304 |
| 126 | KC_release_pos_z_weighted_avg | -0.063441 |
| 127 | FA_release_pos_y_weighted_avg | -0.082745 |
| 128 | FS_release_pos_z_weighted_avg | -0.206523 |


```
In [73]: ▶ 1 # This will rename the index if 'feature' is actually set as the index
2 feature_importances = feature_importances.set_index('feature') # Make
3 feature_importances = feature_importances.rename(index={
4     'season': 'Season',
5     'IN_release_pos_x_weighted_avg': 'Intentional Walk Release Positio
6     'FF_release_pos_y_weighted_avg': 'Fastball Release Position Y-dime
7     'CU_release_pos_x_weighted_avg': 'Curveball Release Position X-dim
8     'FF_release_pos_z_weighted_avg': 'Fastball Release Position Z-dime
9     'CU_release_pos_z_weighted_avg': 'Curveball Release Position Z-dim
10    'CS_release_speed_weighted_avg': 'Slow Curve Release Speed avg.',
11    'FF_release_speed_weighted_avg': 'Fastball Release Speed avg.',
12    'CU_release_speed_weighted_avg': 'Curveball Release Speed avg.',
13    'ST_vy0_weighted_avg': 'Sweeper Velocity Y-dimension avg.',
14    'CH_vz0_weighted_avg': 'Changeup Velocity Z-dimension avg.',
15    'FF_vy0_weighted_avg': 'Fastball Velocity Y-dimension avg.',
16    'ST_release_speed_weighted_avg': 'Sweeper Release Speed avg.',
17    'CS_vy0_weighted_avg': 'Slow Curve Velocity Y-dimension avg.',
18    'PO_release_pos_x_weighted_avg': 'Pick-off Release Position X-dime
19    'ST_release_pos_z_weighted_avg': 'Sweeper Release Position Z-dimer
20    'ST_release_pos_y_weighted_avg': 'Sweeper Release Position Y-dimer
21    'KN_release_pos_y_weighted_avg': 'Knuckleball Release Position Y-d
22    'FC_release_pos_y_weighted_avg': 'Cutter Release Position Y-dimens
23    'KC_release_pos_z_weighted_avg': 'Knucklecurve Release Position Z-
24    'FA_release_pos_y_weighted_avg': '4-seam Fastball Release Positio
25    'FS_release_pos_z_weighted_avg': 'Split-finger Release Position Z-
26
27 })
```

In [74]: 1 feature_importances

Out[74]:

| | importance |
|--|------------|
| feature | |
| Age | 0.350397 |
| Season | 0.204704 |
| Intentional Walk Release Position X-dimension avg. | 0.106614 |
| Fastball Release Position Y-dimension avg. | 0.056850 |
| Curveball Release Position X-dimension avg. | 0.028929 |
| Fastball Release Position Z-dimension avg. | 0.015752 |
| Slow Curve Release Speed avg. | 0.007832 |
| Fastball Release Speed avg. | 0.006973 |
| Sweeper Velocity Y-dimension avg. | 0.005001 |
| Changeup Velocity Z-dimension avg. | 0.004601 |
| Fastball Velocity Y-dimension avg. | -0.012434 |
| Sweeper Release Speed avg. | -0.013466 |
| Throws | -0.018221 |
| Slow Curve Velocity Y-dimension avg. | -0.018479 |
| Pick-off Release Position X-dimension avg. | -0.020661 |
| Sweeper Release Position Z-dimension avg. | -0.024097 |
| Sweeper Release Position Y-dimension avg. | -0.029266 |
| Knuckleball Release Position Y-dimension avg. | -0.033056 |
| Cutter Release Position Y-dimension avg. | -0.036304 |
| Knucklecurve Release Position Z-dimension avg. | -0.063441 |
| 4-seam Fastball Release Position Y-dimension avg. | -0.082745 |
| Split-finger Release Position Z-dimension avg. | -0.206523 |

```
In [75]: 1 sns.barplot(x='importance', y='feature', data=feature_importances, hue=
2 plt.title('Feature Importances')
3 plt.show())
```



This shows the features that have the most impact in predicting 1.0 surgery (positive and negative)

Decision Tree Classifier, baseline model.

```
In [ ]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [ ]: 1 tree_clf = DecisionTreeClassifier(criterion='gini', max_depth=5)
2 tree_clf.fit(X_train, y_train)
```

Out[31]: DecisionTreeClassifier(max_depth=5)


In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```

In [ ]: ▶ 1 def plot_feature_importances(model, n_top_features=20):
2     importances = model.feature_importances_
3     indices = np.argsort(importances)[-n_top_features:]
4     plt.figure(figsize=(10,10))
5     plt.title('Top N Feature Importances')
6     plt.barh(range(n_top_features), importances[indices], align='center')
7     plt.yticks(range(n_top_features), [X_train.columns[i] for i in indices])
8     plt.xlabel('Feature importance')
9     plt.ylabel('Feature')
10    plt.ylim(-1, n_top_features)
11
12    plot_feature_importances(tree_clf, n_top_features=20)
13    plt.show()

```




```

In [ ]: ▶ 1 pred = tree_clf.predict(X_test)

```

```
In [ ]: 1 pred = tree_clf.predict(X_test)
2
3 ConfusionMatrixDisplay.from_predictions(y_test, pred)
4 plt.title('pivoted_df Decision Tree Classifier')
5 plt.show()
6 print(classification_report(y_test, pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.77 | 0.98 | 0.86 | 827 |
| 1.0 | 0.70 | 0.11 | 0.19 | 280 |
| accuracy | | | 0.76 | 1107 |
| macro avg | 0.74 | 0.55 | 0.53 | 1107 |
| weighted avg | 0.75 | 0.76 | 0.69 | 1107 |

Terrible for TP and FP. Need to adjust. Features are interesting. Mostly fastball, curveball, some slider and split-finger.

```
In [ ]: 1 y = pivoted_df['Surgery']
2 X = pivoted_df.drop('Surgery', axis=1)
```

```
In [ ]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
```

```
In [ ]: 1 param_grid = {
2     'criterion': ['gini', 'entropy'],
3     'max_depth': [5, 10, 15, 20],
4     'min_samples_split': [2, 5, 10],
5     'min_samples_leaf': [1, 2, 4],
6     'class_weight': ['balanced', {0:1, 1:2}, {0:1, 1:3}]
7 }
8
9 tree_clf = DecisionTreeClassifier()
10 scorer = make_scorer(recall_score)
11 grid_search = GridSearchCV(estimator=tree_clf, param_grid=param_grid,
12 grid_search.fit(X_train, y_train)
13
14 print("Best parameters:", grid_search.best_params_)
15 print("Best score:", grid_search.best_score_)
16
17 best_tree = grid_search.best_estimator_
18 y_pred = best_tree.predict(X_test)
19 print("Test recall score:", recall_score(y_test, y_pred))
```

Best parameters: {'class_weight': 'balanced', 'criterion': 'gini', 'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 10}
Best score: 0.6524852362204725
Test recall score: 0.7607142857142857

```
In [ ]: 1 best_tree.fit(X_train, y_train)
```


Out[38]: DecisionTreeClassifier(class_weight='balanced', max_depth=5, min_samples_leaf=2,
min_samples_split=10)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [ ]: 1 pred = best_tree.predict(X_test)
```

```
In [ ]: 1 ConfusionMatrixDisplay.from_predictions(y_test, pred)
2 plt.title('Grid Search Recall Decision Tree Classifier')
3 plt.show()
4 print(classification_report(y_test, pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.87 | 0.52 | 0.65 | 827 |
| 1.0 | 0.35 | 0.76 | 0.48 | 280 |
| accuracy | | | 0.58 | 1107 |
| macro avg | 0.61 | 0.64 | 0.56 | 1107 |
| weighted avg | 0.74 | 0.58 | 0.61 | 1107 |

The Logistic Regression model with adjusted class weights performed the best.

```
In [ ]: 1
```