

```

In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 %matplotlib inline
        5 import seaborn as sns
        6 from sklearn.utils.class_weight import compute_class_weight
        7 from sklearn.preprocessing import StandardScaler
        8 from sklearn.linear_model import LogisticRegression
        9 from sklearn.tree import DecisionTreeClassifier
       10 from sklearn.ensemble import RandomForestClassifier
       11 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
       12 from sklearn.metrics import accuracy_score, recall_score, precision_score
       13 from sklearn.metrics import ConfusionMatrixDisplay
       14 from sklearn.metrics import classification_report
       15 from sklearn.pipeline import Pipeline
       16 from imblearn.pipeline import Pipeline as ImbPipeline
       17 from sklearn.decomposition import PCA
       18 from imblearn.over_sampling import SMOTE, BorderlineSMOTE
       19 from google.colab import files
       20 uploaded = files.upload()

```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session.  
Please rerun this cell to enable.

Saving pivoted\_df.csv to pivoted\_df.csv

Load the dataframe in, inspect the data.

```

In [5]: 1 pivoted_df = pd.read_csv('~/.Documents/Flatiron/Project_5_/data/pivoted_

```

```

In [6]: 1 pivoted_df.head()

```

Out[6]:

	season	Age	Throws	Surgery	AB_release_speed_weighted_avg	CH_release_speed_weight
0	2008	37.0	1	0.0	0.0	8.0
1	2009	38.0	1	0.0	0.0	8.0
2	2010	39.0	1	0.0	0.0	8.0
3	2011	40.0	1	0.0	0.0	8.0
4	2012	41.0	1	0.0	0.0	8.0

5 rows × 130 columns

```

In [7]: 1 pivoted_df.shape

```

Out[7]: (3688, 130)

In [8]: 1 pivoted\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3688 entries, 0 to 3687
Columns: 130 entries, season to SV_vz0_weighted_avg
dtypes: float64(128), int64(2)
memory usage: 3.7 MB
```

In [9]: 1 pivoted\_df['Surgery'].value\_counts()

```
Out[9]: 0.0    2772
        1.0     916
        Name: Surgery, dtype: int64
```

Time to start modeling! Split target and features and make a baseline model.

In [10]: 1 y = pivoted\_df['Surgery']  
2 X = pivoted\_df.drop('Surgery', axis=1)

In [11]: 1 X\_train, X\_test, y\_train, y\_test = train\_test\_split(X, y, test\_size=0.

```

In [12]: 1 #Make a pipeline to simplify process
2 logreg_pipeline = Pipeline([
3     ('scale', StandardScaler()),
4     ('logreg', LogisticRegression(solver='liblinear'))
5 ])
6
7 # Define parameter grid to search
8 param_grid = {
9     'logreg_C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
10    'logreg_penalty': ['l1', 'l2'] # Norm used in the penalization
11 }
12
13 # Initialize GridSearchCV with pipeline, parameter grid, and scoring method
14 grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='roc_auc')
15
16 # Assuming X_train and y_train are already defined
17 grid_search.fit(X_train, y_train)
18
19 # Best parameters found
20 print("Best parameters: ", grid_search.best_params_)
21
22 # Best cross-validation score
23 print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
24
25 # Test set score using the best parameters
26 print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))

```

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

Best parameters: {'logreg\_C': 10, 'logreg\_penalty': 'l1'}

Best cross-validation score: 0.76

Test set score: 0.77

```
In [13]: 1 logreg_pipeline = Pipeline([
2         ('scale', StandardScaler()),
3         ('logreg', LogisticRegression(penalty='l1', C=10.0, solver='liblinear')),
4     ])
```

```
In [14]: 1 logreg_pipeline.fit(X_train, y_train)
```

```
Out[14]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(C=10.0, penalty='l1', solver='liblinear'))])
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

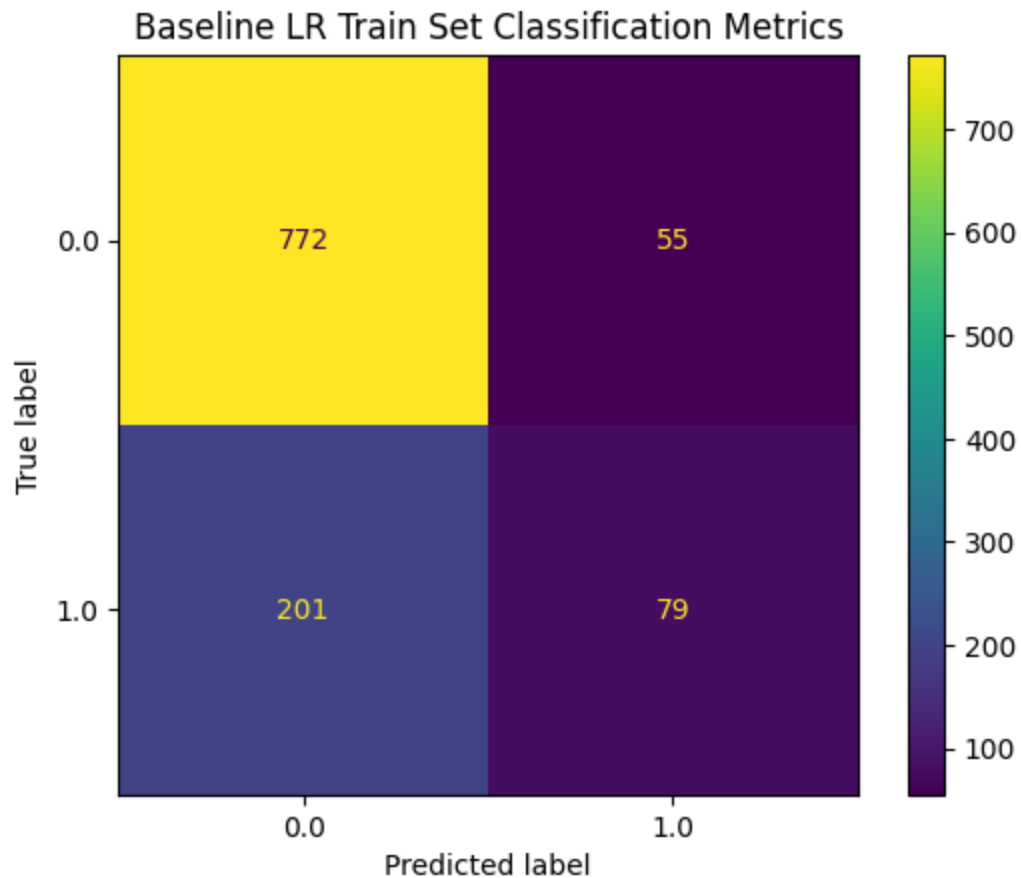
**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [15]: 1 logreg_pipeline.score(X_test, y_test)
```

```
Out[15]: 0.7687443541102078
```

```
In [16]: 1 y_pred = logreg_pipeline.predict(X_test)
```

```
In [17]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Baseline LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0.0	0.79	0.93	0.86	827
1.0	0.59	0.28	0.38	280
accuracy			0.77	1107
macro avg	0.69	0.61	0.62	1107
weighted avg	0.74	0.77	0.74	1107

Dataset is imbalanced, need to adjust. Should also focus on Recall score since this is a medical issue (better to have False Positive than True Negative!)

```
In [18]: 1 y = pivoted_df['Surgery']
2 X = pivoted_df.drop('Surgery', axis=1)
```

```
In [19]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
```

```

In [24]: 1 # Set up pipeline
2 weight_logreg_pipeline = Pipeline([
3     ('scale', StandardScaler()),
4     ('logreg', LogisticRegression(solver='liblinear'))
5 ])
6
7 # Define the parameter grid to search over, including class weights
8 param_grid = {
9     'logreg__C': [0.01, 0.1, 1, 10],
10    'logreg__penalty': ['l1', 'l2'],
11    'logreg__class_weight': [None, 'balanced', {0: 1, 1: 2}, {0: 1, 1:
12    'logreg__max_iter': [5000],
13    'logreg__tol': [0.01]
14 }
15
16 # Create a scoring function that focuses on recall for the positive cl
17 recall_scorer = make_scorer(recall_score, pos_label=1)
18
19 # Initialize GridSearch with pipeline, param grid, and recall
20 grid_search = GridSearchCV(weight_logreg_pipeline, param_grid, cv=5, s
21
22 # Fit the grid search to the data
23 grid_search.fit(X_train, y_train)
24
25 # Print the best parameters found and the best recall score
26 print("Best parameters: ", grid_search.best_params_)
27 print("Best cross-validation recall score: {:.2f}".format(grid_search.
28
29 # Evaluate the best model on the test set
30 best_model = grid_search.best_estimator_
31 y_pred = best_model.predict(X_test)
32 print("Test set recall score: {:.2f}".format(recall_score(y_test, y_pr

```

```

Best parameters: {'logreg__C': 0.01, 'logreg__class_weight': {0: 1, 1:
5}, 'logreg__max_iter': 5000, 'logreg__penalty': 'l1', 'logreg__tol': 0.0
1}
Best cross-validation recall score: 0.87
Test set recall score: 0.86

```

```

In [25]: 1 best_model.fit(X_train, y_train)

```

```

Out[25]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(C=0.01, class_weight={0: 1, 1: 5},
                                                max_iter=5000, penalty='l1',
                                                solver='liblinear', tol=0.01))])

```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [26]: 1 best_model.score(X_test, y_test)
```

```
Out[26]: 0.5149051490514905
```

```
In [28]: 1 y_pred = best_model.predict(X_test)
```

```
In [29]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Weighted LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0.0	0.89	0.40	0.55	827
1.0	0.33	0.85	0.47	280
accuracy			0.51	1107
macro avg	0.61	0.63	0.51	1107
weighted avg	0.75	0.51	0.53	1107

Much better model. False Negatives is low, other classes much higher.

Decision Tree Classifier, baseline model.

```
In [30]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
```


```
In [31]: 1 tree_clf = DecisionTreeClassifier(criterion='gini', max_depth=5)
2 tree_clf.fit(X_train, y_train)
```

Out[31]: DecisionTreeClassifier(max\_depth=5)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

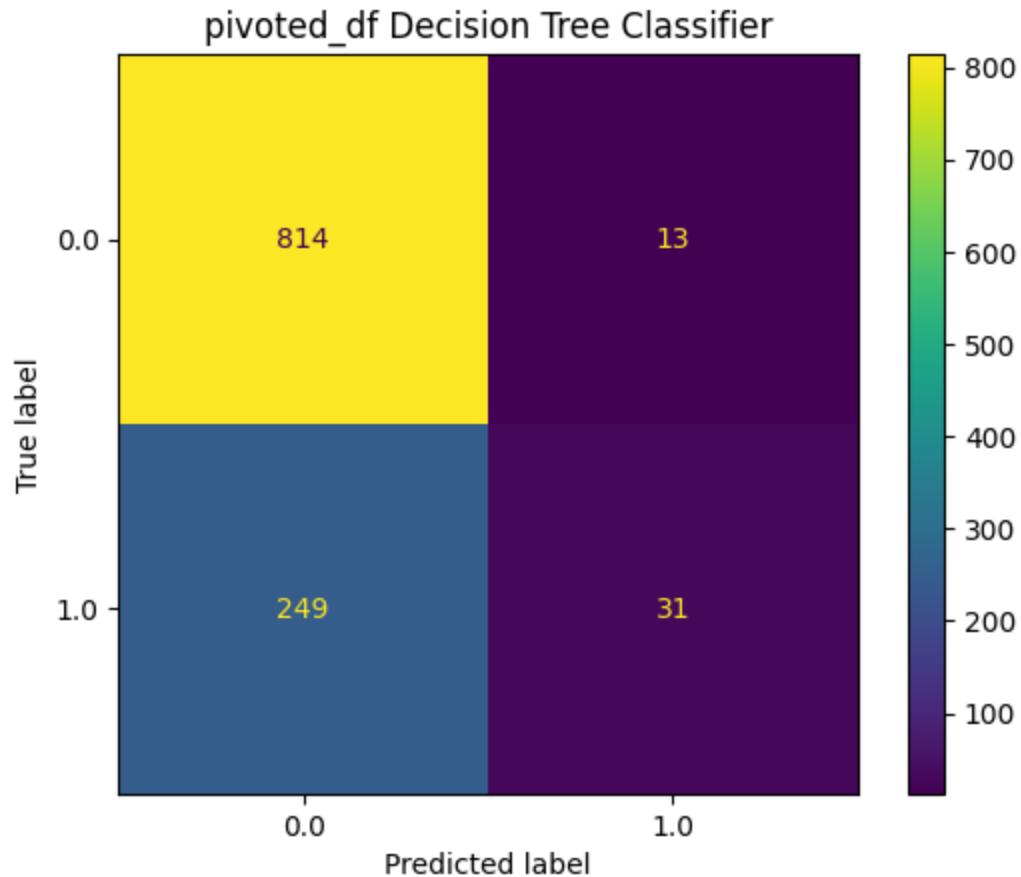
```
In [32]: 1 def plot_feature_importances(model, n_top_features=20):
2     importances = model.feature_importances_
3     indices = np.argsort(importances)[-n_top_features:]
4     plt.figure(figsize=(10,10))
5     plt.title('Top N Feature Importances')
6     plt.barh(range(n_top_features), importances[indices], align='center')
7     plt.yticks(range(n_top_features), [X_train.columns[i] for i in indices])
8     plt.xlabel('Feature importance')
9     plt.ylabel('Feature')
10    plt.ylim(-1, n_top_features)
11
12    plot_feature_importances(tree_clf, n_top_features=20)
13    plt.show()
```





In [33]: `1 pred = tree_clf.predict(X_test)`

In [34]: `1 pred = tree_clf.predict(X_test)  
2  
3 ConfusionMatrixDisplay.from_predictions(y_test, pred)  
4 plt.title('pivoted_df Decision Tree Classifier')  
5 plt.show()  
6 print(classification_report(y_test, pred))`



	precision	recall	f1-score	support
0.0	0.77	0.98	0.86	827
1.0	0.70	0.11	0.19	280
accuracy			0.76	1107
macro avg	0.74	0.55	0.53	1107
weighted avg	0.75	0.76	0.69	1107

Terrible for TP and FP. Need to adjust. Features are interesting. Mostly fastball, curveball, some slider and split-finger.

In [35]: `1 y = pivoted_df['Surgery']  
2 X = pivoted_df.drop('Surgery', axis=1)`

```
In [36]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
```

```
In [37]: 1 param_grid = {
2     'criterion': ['gini', 'entropy'],
3     'max_depth': [5, 10, 15, 20],
4     'min_samples_split': [2, 5, 10],
5     'min_samples_leaf': [1, 2, 4],
6     'class_weight': ['balanced', {0:1, 1:2}, {0:1, 1:3}]
7 }
8
9 tree_clf = DecisionTreeClassifier()
10 scorer = make_scorer(recall_score)
11 grid_search = GridSearchCV(estimator=tree_clf, param_grid=param_grid,
12 grid_search.fit(X_train, y_train)
13
14 print("Best parameters:", grid_search.best_params_)
15 print("Best score:", grid_search.best_score_)
16
17 best_tree = grid_search.best_estimator_
18 y_pred = best_tree.predict(X_test)
19 print("Test recall score:", recall_score(y_test, y_pred))
```

```
Best parameters: {'class_weight': 'balanced', 'criterion': 'gini', 'max_d
epth': 5, 'min_samples_leaf': 2, 'min_samples_split': 10}
Best score: 0.6524852362204725
Test recall score: 0.7607142857142857
```

```
In [38]: 1 best_tree.fit(X_train, y_train)
```


```
Out[38]: DecisionTreeClassifier(class_weight='balanced', max_depth=5, min_samples_
leaf=2,
                                min_samples_split=10)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [39]: 1 pred = best_tree.predict(X_test)
```

```
In [40]: 1 ConfusionMatrixDisplay.from_predictions(y_test, pred)
2 plt.title('Grid Search Recall Decision Tree Classifier')
3 plt.show()
4 print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.87	0.52	0.65	827
1.0	0.35	0.76	0.48	280
accuracy			0.58	1107
macro avg	0.61	0.64	0.56	1107
weighted avg	0.74	0.58	0.61	1107

The Logistic Regression model with adjusted class weights performed the best.

```
In [ ]: 1
```