

```
In [1]: ▶ 1 import pybaseball as pyb
2 from pybaseball import statcast, pitching_stats, playerid_lookup, stat
3 import numpy as np
4 import math
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 %matplotlib inline
8 import seaborn as sns
9 import glob
10 import os
11 import re
12 import unicodedata
13 from datetime import datetime
14 from itertools import groupby
15 from operator import itemgetter
16 from sklearn.preprocessing import OneHotEncoder, StandardScaler
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.model_selection import train_test_split, GridSearchCV
19 from sklearn.metrics import accuracy_score, recall_score, precision_sc
20 from sklearn.metrics import ConfusionMatrixDisplay
21 from sklearn.metrics import classification_report
22 from sklearn.pipeline import Pipeline
23 #from imblearn.over_sampling import SMOTE
24 from catboost import CatBoostClassifier
```

C:\Users\johns\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:6
0: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck'
(version '1.3.5' currently installed).
from pandas.core import (


```
In [2]: ▶ 1 complete_100_df = pd.read_csv('~\Documents\Flatiron\Project_5_\data\com
```

In [24]: 1 complete_100_df

Out[24]:

	Name	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	relea
0	adam wainwright	41.0	425794	2023	CH		91
1	adam wainwright	41.0	425794	2023	CS		3
2	adam wainwright	41.0	425794	2023	CU		545
3	adam wainwright	41.0	425794	2023	FC		403
4	adam wainwright	41.0	425794	2023	FF		176
...
21386	jeff samardzija	23.0	502188	2008	FS		99
21387	jeff samardzija	23.0	502188	2008	IN		6
21388	jeff samardzija	23.0	502188	2008	PO		1
21389	jeff samardzija	23.0	502188	2008	SI		83
21390	jeff samardzija	23.0	502188	2008	SL		45

21391 rows × 16 columns



In [26]: 1 complete_100_df['Surgery'].value_counts()

Out[26]: Surgery
 0.0 16389
 1.0 5002
 Name: count, dtype: int64

In [23]: ▶

1

wainwright_df = complete_100_df[complete_100_df['Name'] == 'adam wainw


2

wainwright_df

Out[23]:

	Name	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	relea
0	adam wainwright	41.0	425794	2023	CH		91
1	adam wainwright	41.0	425794	2023	CS		3
2	adam wainwright	41.0	425794	2023	CU		545
3	adam wainwright	41.0	425794	2023	FC		403
4	adam wainwright	41.0	425794	2023	FF		176
...
20627	adam wainwright	26.0	425794	2008	FC		395
20628	adam wainwright	26.0	425794	2008	FF		101
20629	adam wainwright	26.0	425794	2008	IN		4
20630	adam wainwright	26.0	425794	2008	PO		2
20631	adam wainwright	26.0	425794	2008	SI		879

95 rows × 16 columns



```
In [13]: 1 # Correcting the approach to ensure 'Surgery' is 1.0 from the surgery
2 complete_100_df['Surgery'] = complete_100_df.apply(lambda row: 1.0 if
3
4 complete_100_df
```

Out[13]:

	Name	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	relea
0	adam wainwright	41.0	425794	2023	CH		91
1	adam wainwright	41.0	425794	2023	CS		3
2	adam wainwright	41.0	425794	2023	CU		545
3	adam wainwright	41.0	425794	2023	FC		403
4	adam wainwright	41.0	425794	2023	FF		176
...
21386	jeff samardzija	23.0	502188	2008	FS		99
21387	jeff samardzija	23.0	502188	2008	IN		6
21388	jeff samardzija	23.0	502188	2008	PO		1
21389	jeff samardzija	23.0	502188	2008	SI		83
21390	jeff samardzija	23.0	502188	2008	SL		45

21391 rows × 16 columns

This ensures that we are only counting for 'Surgery' for the year of surgery, and all years going forward.

```
In [17]: 1 filtered_rows = complete_100_df[(complete_100_df['Surgery'] == 1.0) &
2 filtered_rows
```

Out[17]:

	Name	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	release_speed_
--	------	-----	---------	--------	------------	----------------------------------	----------------

```
In [19]: 1 pd.set_option('display.max_rows', None)
```

```
In [22]: 1 pd.reset_option('display.max_rows')
```

```
In [27]: 1 complete_100_df['TJ Surgery Year'].fillna(0.0, inplace=True)
```

C:\Users\johns\AppData\Local\Temp\ipykernel_8568\550959113.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
complete_100_df['TJ Surgery Year'].fillna(0.0, inplace=True)
```

```
In [7]: ▶ 1 """
2 # Fill NaN in 'TJ Surgery Year' with 0.0
3 complete_100_df['TJ Surgery Year'].fillna(0.0, inplace=True)
4
5 # Update 'Surgery' based on 'season' and 'TJ Surgery Year'
6 complete_100_df['Surgery'] = complete_100_df.apply(
7     lambda row: 1.0 if row['season'] >= row['TJ Surgery Year'] and row
8 )
9
10 # Verify the changes
11 print(complete_100_df[['Name', 'season', 'TJ Surgery Year', 'Surgery']]
12
13 # Check the value counts again
14 print(complete_100_df['Surgery'].value_counts())
15 """
```

C:\Users\johns\AppData\Local\Temp\ipykernel_32520\2928425957.py:2: Future Warning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
complete_100_df['TJ Surgery Year'].fillna(0.0, inplace=True)
```

	Name	season	TJ Surgery Year	Surgery
0	adam wainwright	2023	2011.0	1.0
1	adam wainwright	2023	2011.0	1.0
2	adam wainwright	2023	2011.0	1.0
3	adam wainwright	2023	2011.0	1.0
4	adam wainwright	2023	2011.0	1.0

Surgery

0.0	16389
1.0	5002

Name: count, dtype: int64

In [28]: 1 complete_100_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 21391 entries, 0 to 21390
Data columns (total 16 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   Name                                           21391 non-null  object
1   Age                                           21391 non-null  float64
2   pitcher                                       21391 non-null  int64
3   season                                       21391 non-null  int64
4   pitch_type                                    21391 non-null  object
5   season_total_count_by_pitch_type            21391 non-null  int64
6   release_speed_weighted_avg                  21391 non-null  float64
7   release_pos_x_weighted_avg                  21391 non-null  float64
8   release_pos_y_weighted_avg                  21391 non-null  float64
9   release_pos_z_weighted_avg                  21391 non-null  float64
10  vx0_weighted_avg                             21391 non-null  float64
11  vy0_weighted_avg                             21391 non-null  float64
12  vz0_weighted_avg                             21391 non-null  float64
13  Throws                                         21391 non-null  int64
14  Surgery                                       21391 non-null  float64
15  TJ Surgery Year                             21391 non-null  float64
dtypes: float64(10), int64(4), object(2)
memory usage: 2.8+ MB
```

Can probably drop 'Name' and 'TJ Surgery Year' columns

In [66]: 1 funky_df = complete_100_df.drop(columns=['Name', 'TJ Surgery Year'])

In [67]: 1 funky_df

Out[67]:

	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	release_speed_w
0	41.0	425794	2023	CH		91
1	41.0	425794	2023	CS		3
2	41.0	425794	2023	CU		545
3	41.0	425794	2023	FC		403
4	41.0	425794	2023	FF		176
...
21386	23.0	502188	2008	FS		99
21387	23.0	502188	2008	IN		6
21388	23.0	502188	2008	PO		1
21389	23.0	502188	2008	SI		83
21390	23.0	502188	2008	SL		45

21391 rows × 14 columns

In [68]: 1 funky_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 21391 entries, 0 to 21390
Data columns (total 14 columns):
 #   Column                                          Non-Null Count  Dtype
---  -
 0   Age                                           21391 non-null  float64
 1   pitcher                                       21391 non-null  int64
 2   season                                       21391 non-null  int64
 3   pitch_type                                    21391 non-null  object
 4   season_total_count_by_pitch_type            21391 non-null  int64
 5   release_speed_weighted_avg                  21391 non-null  float64
 6   release_pos_x_weighted_avg                  21391 non-null  float64
 7   release_pos_y_weighted_avg                  21391 non-null  float64
 8   release_pos_z_weighted_avg                  21391 non-null  float64
 9   vx0_weighted_avg                            21391 non-null  float64
10   vy0_weighted_avg                            21391 non-null  float64
11   vz0_weighted_avg                            21391 non-null  float64
12   Throws                                       21391 non-null  int64
13   Surgery                                       21391 non-null  float64
dtypes: float64(9), int64(4), object(1)
memory usage: 2.4+ MB
```

In [69]: 1 funky_df['pitch_type'].value_counts()

```
Out[69]: pitch_type
FF      3641
CH      3356
SI      3286
SL      2775
CU      2670
FC      1625
IN      1605
PO      1045
KC       525
FS       390
ST       129
FA       124
EP        90
CS        52
SV        29
KN        25
AB        14
SC         10
Name: count, dtype: int64
```


In [70]: 1 funky_df['release_pos_y_weighted_avg'].value_counts()

Out[70]: release_pos_y_weighted_avg
 54.500000 14019
 54.580000 7
 54.780000 6
 54.360000 6
 54.160000 6
 ...
 54.247143 1
 54.289964 1
 54.083586 1
 54.099168 1
 55.031195 1
 Name: count, Length: 6936, dtype: int64

In [58]: 1 fa_rows = funky_df[funky_df['pitch_type'] == 'FA']

In [59]: 1 fa_rows

Out[59]:

	Age	pitcher	season	pitch_type	season_total_count_by_pitch_type	release_speed_w
67	35.0	477132	2023	FA		1
168	33.0	543101	2023	FA		1
196	32.0	543475	2023	FA		1
1869	37.0	425844	2021	FA		9
2892	36.0	425844	2020	FA		57
...
20964	23.0	444836	2008	FA		2
21019	27.0	446454	2008	FA		7
21211	25.0	456043	2008	FA		20
21219	22.0	456501	2008	FA		1
21228	26.0	456589	2008	FA		6

124 rows × 14 columns

Try condensing pitch_type before the pivot and compare.

In [60]: 1 condensed_pitch_type_df = funky_df

```
In [61]: 1 pitch_type_mapping = {
2         'FF': 'FB', 'SI': 'FB', 'FC': 'FB', 'FA': 'FB',
3         'CH': 'OS', 'FS': 'OS', 'FO': 'OS', 'SC': 'OS', 'PO': 'OS',
4         'CU': 'BB', 'KC': 'BB', 'CS': 'BB',
5         'SL': 'SB', 'ST': 'SB', 'SV': 'SB', 'KN': 'SB',
6         'EP': 'OT', 'AB': 'OT', 'IN': 'OT'
7     }
8
9 condensed_pitch_type_df['pitch_type_group'] = condensed_pitch_type_df[
```

```
In [63]: 1 grouped_df = condensed_pitch_type_df.groupby(['Age', 'pitcher', 'season',
2         season_total_count_by_pitch_type=('season_total_count_by_pitch_type', 'mean'),
3         release_speed_weighted_avg=('release_speed_weighted_avg', 'mean'),
4         release_pos_x_weighted_avg=('release_pos_x_weighted_avg', 'mean'),
5         release_pos_y_weighted_avg=('release_pos_y_weighted_avg', 'mean'),
6         release_pos_z_weighted_avg=('release_pos_z_weighted_avg', 'mean'),
7         vx0_weighted_avg=('vx0_weighted_avg', 'mean'),
8         vy0_weighted_avg=('vy0_weighted_avg', 'mean'),
9         vz0_weighted_avg=('vz0_weighted_avg', 'mean'),
10        Throws=('Throws', 'first'), # Assuming Throws doesn't change with
11        Surgery=('Surgery', 'first') # Assuming Surgery doesn't change with
12    ).reset_index()
13
```

```
In [64]: 1 grouped_df
```

```
Out[64]:
```

	Age	pitcher	season	pitch_type_group	season_total_count_by_pitch_type	release_sp
0	19.0	518516	2009	BB		42
1	19.0	518516	2009	FB		98
2	19.0	518516	2009	OS		12
3	19.0	518516	2009	OT		4
4	19.0	605164	2012	BB		1
...
14707	47.0	119469	2010	OS		306
14708	49.0	119469	2012	BB		90
14709	49.0	119469	2012	FB		628
14710	49.0	119469	2012	OS		301
14711	49.0	119469	2012	OT		8

14712 rows × 14 columns

```
In [65]: 1 grouped_df['release_pos_y_weighted_avg'].value_counts()
```

```
Out[65]: release_pos_y_weighted_avg
54.500000    9540
54.160000      3
53.970000      3
53.980000      3
54.110000      3
...
54.281383      1
54.396867      1
53.962602      1
54.059634      1
55.724512      1
Name: count, Length: 5138, dtype: int64
```

```

In [72]: 1 def pivot_metrics(df, index_cols, pivot_col, value_cols):
2         """
3         Pivot the DataFrame for the specified pivot column.
4         :param df: DataFrame to pivot.
5         :param index_cols: List of columns to use as the index.
6         :param pivot_col: Column to pivot on.
7         :param value_cols: Columns whose values are to be spread across pivot.
8         :return: Pivoted DataFrame.
9         """
10        pivoted_dfs = []
11        for value_col in value_cols:
12            # Pivot each metric column separately and rename to include the metric name
13            pivoted_df = df.pivot_table(index=index_cols, columns=pivot_col, values=value_col)
14            pivoted_df.columns = [f"{col}_{value_col}" if col not in index_cols else col for col in pivoted_df.columns]
15            pivoted_dfs.append(pivoted_df)
16
17        # Merge all the pivoted metric DataFrames on the index columns
18        from functools import reduce
19        final_df = reduce(lambda left, right: pd.merge(left, right, on=index_cols, how='outer'), pivoted_dfs)
20        return final_df
21
22        # Define the base columns and the metrics you want to pivot
23        index_cols = ['pitcher', 'season', 'Age', 'Throws', 'Surgery']
24        pivot_col = 'pitch_type_group'
25        value_cols = ['release_speed_weighted_avg', 'release_pos_x_weighted_avg', 'release_pos_y_weighted_avg']
26
27        # Pivot the DataFrame
28        cond_pivoted_df = pivot_metrics(grouped_df, index_cols, pivot_col, value_cols)
29
30        cond_pivoted_df.head()

```

Out[72]:

	pitcher	season	Age	Throws	Surgery	BB_release_speed_weighted_avg	FB_release_speed_weighted_avg
0	110683	2008	37.0	1	0.0	75.425843	75.425843
1	110683	2009	38.0	1	0.0	78.181818	78.181818
2	110683	2010	39.0	1	0.0	74.666667	74.666667
3	110683	2011	40.0	1	0.0	76.885714	76.885714
4	110683	2012	41.0	1	0.0	76.427273	76.427273

5 rows × 40 columns

In [73]: `cond_pivoted_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3688 entries, 0 to 3687
Data columns (total 40 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   pitcher                                                                3688 non-null   int64
1   season                                                                3688 non-null   int64
2   Age                                                                    3688 non-null   float64
3   Throws                                                                3688 non-null   int64
4   Surgery                                                                3688 non-null   float64
5   BB_release_speed_weighted_avg    3059 non-null   float64
6   FB_release_speed_weighted_avg    3688 non-null   float64
7   OS_release_speed_weighted_avg    3562 non-null   float64
8   OT_release_speed_weighted_avg    1610 non-null   float64
9   SB_release_speed_weighted_avg    2793 non-null   float64
10  BB_release_pos_x_weighted_avg    3059 non-null   float64
11  FB_release_pos_x_weighted_avg    3688 non-null   float64
12  OS_release_pos_x_weighted_avg    3562 non-null   float64
13  OT_release_pos_x_weighted_avg    1610 non-null   float64
14  SB_release_pos_x_weighted_avg    2793 non-null   float64
15  BB_release_pos_y_weighted_avg    3059 non-null   float64
16  FB_release_pos_y_weighted_avg    3688 non-null   float64
17  OS_release_pos_y_weighted_avg    3562 non-null   float64
18  OT_release_pos_y_weighted_avg    1610 non-null   float64
19  SB_release_pos_y_weighted_avg    2793 non-null   float64
20  BB_release_pos_z_weighted_avg    3059 non-null   float64
21  FB_release_pos_z_weighted_avg    3688 non-null   float64
22  OS_release_pos_z_weighted_avg    3562 non-null   float64
23  OT_release_pos_z_weighted_avg    1610 non-null   float64
24  SB_release_pos_z_weighted_avg    2793 non-null   float64
25  BB_vx0_weighted_avg            3059 non-null   float64
26  FB_vx0_weighted_avg            3688 non-null   float64
27  OS_vx0_weighted_avg            3562 non-null   float64
28  OT_vx0_weighted_avg            1610 non-null   float64
29  SB_vx0_weighted_avg            2793 non-null   float64
30  BB_vy0_weighted_avg            3059 non-null   float64
31  FB_vy0_weighted_avg            3688 non-null   float64
32  OS_vy0_weighted_avg            3562 non-null   float64
33  OT_vy0_weighted_avg            1610 non-null   float64
34  SB_vy0_weighted_avg            2793 non-null   float64
35  BB_vz0_weighted_avg            3059 non-null   float64
36  FB_vz0_weighted_avg            3688 non-null   float64
37  OS_vz0_weighted_avg            3562 non-null   float64
38  OT_vz0_weighted_avg            1610 non-null   float64
39  SB_vz0_weighted_avg            2793 non-null   float64
dtypes: float64(37), int64(3)
memory usage: 1.1 MB
```

This condensed DF has 40 columns compared to before where I had 130 columns.

```
In [74]: 1 cond_pivoted_df.fillna(0.0, inplace=True)

In [75]: 1 cond_pivoted_df.to_csv('data/cond_pivoted_df.csv')

In [76]: 1 cond_groovy_df = cond_pivoted_df

In [77]: 1 cond_groovy_df.drop(columns=['pitcher'], inplace=True)

In [78]: 1 y = cond_groovy_df['Surgery']
          2 X = cond_groovy_df.drop('Surgery', axis=1)

In [79]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
```

```

In [80]: ► 1 logreg_pipeline = Pipeline([
2           ('scale', StandardScaler()),
3           ('logreg', LogisticRegression(solver='liblinear'))
4       ])
5
6       # Define the parameter grid to search over
7       param_grid = {
8           'logreg__C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
9           'logreg__penalty': ['l1', 'l2'] # Norm used in the penalization
10      }
11
12      # Initialize GridSearchCV with the pipeline, parameter grid, and desired cv
13      grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='roc_auc')
14
15      # Assuming X_train and y_train are already defined
16      grid_search.fit(X_train, y_train)
17
18      # Best parameters found
19      print("Best parameters: ", grid_search.best_params_)
20
21      # Best cross-validation score
22      print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
23
24      # Test set score using the best parameters
25      print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))

```

C:\Users\johns\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

C:\Users\johns\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

C:\Users\johns\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

C:\Users\johns\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

C:\Users\johns\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(

Best parameters: {'logreg__C': 1, 'logreg__penalty': 'l1'}

Best cross-validation score: 0.75

Test set score: 0.75

```
In [81]: 1 logreg_pipeline = Pipeline([
2         ('scale', StandardScaler()),
3         ('logreg', LogisticRegression(penalty='l1', C=1.0, solver='liblinear',
4         ])
```

```
In [82]: 1 logreg_pipeline.fit(X_train, y_train)
```

```
Out[82]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(penalty='l1', solver='liblinear'))])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [83]: 1 logreg_pipeline.score(X_test, y_test)
```

```
Out[83]: 0.7497741644083108
```

```
In [84]: 1 y_pred = logreg_pipeline.predict(X_test)
```



```
In [85]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Baseline LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0.0	0.77	0.96	0.85	827
1.0	0.52	0.14	0.22	280
accuracy			0.75	1107
macro avg	0.64	0.55	0.53	1107
weighted avg	0.70	0.75	0.69	1107

Recall score was even worse for this. Will try with SMOTE on google colab

```

In [31]: ▶ 1 def pivot_metrics(df, index_cols, pivot_col, value_cols):
2         """
3         Pivot the DataFrame for the specified pivot column.
4         :param df: DataFrame to pivot.
5         :param index_cols: List of columns to use as the index.
6         :param pivot_col: Column to pivot on.
7         :param value_cols: Columns whose values are to be spread across pivot.
8         :return: Pivoted DataFrame.
9         """
10        pivoted_dfs = []
11        for value_col in value_cols:
12            # Pivot each metric column separately and rename to include the metric name
13            pivoted_df = df.pivot_table(index=index_cols, columns=pivot_col, values=value_col)
14            pivoted_df.columns = [f"{col}_{value_col}" if col not in index_cols else col for col in pivoted_df.columns]
15            pivoted_dfs.append(pivoted_df)
16
17        # Merge all the pivoted metric DataFrames on the index columns
18        from functools import reduce
19        final_df = reduce(lambda left, right: pd.merge(left, right, on=index_cols, how='outer'), pivoted_dfs)
20        return final_df
21
22        # Define the base columns and the metrics you want to pivot
23        index_cols = ['pitcher', 'season', 'Age', 'Throws', 'Surgery']
24        pivot_col = 'pitch_type'
25        value_cols = ['release_speed_weighted_avg', 'release_pos_x_weighted_avg', 'release_pos_y_weighted_avg']
26
27        # Pivot the DataFrame
28        pivoted_df = pivot_metrics(funky_df, index_cols, pivot_col, value_cols)
29
30        pivoted_df.head()

```

Out[31]:

	pitcher	season	Age	Throws	Surgery	AB_release_speed_weighted_avg	CH_release_speed_weighted_avg
0	110683	2008	37.0	1	0.0	NaN	NaN
1	110683	2009	38.0	1	0.0	NaN	NaN
2	110683	2010	39.0	1	0.0	NaN	NaN
3	110683	2011	40.0	1	0.0	NaN	NaN
4	110683	2012	41.0	1	0.0	NaN	NaN

5 rows × 131 columns

In [36]: 1 pivoted_df

Out[36]:

	pitcher	season	Age	Throws	Surgery	AB_release_speed_weighted_avg	CH_release_
0	110683	2008	37.0	1	0.0		0.0
1	110683	2009	38.0	1	0.0		0.0
2	110683	2010	39.0	1	0.0		0.0
3	110683	2011	40.0	1	0.0		0.0
4	110683	2012	41.0	1	0.0		0.0
...
3683	672578	2022	25.0	1	0.0		0.0
3684	672578	2023	26.0	1	0.0		0.0
3685	680686	2021	23.0	1	0.0		0.0
3686	680686	2022	24.0	1	0.0		0.0
3687	680686	2023	25.0	1	0.0		0.0

3688 rows × 131 columns

In [35]: 1 pivoted_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3688 entries, 0 to 3687
Columns: 131 entries, pitcher to SV_vz0_weighted_avg
dtypes: float64(128), int64(3)
memory usage: 3.7 MB
```

In [34]: 1 pivoted_df.fillna(0.0, inplace=True)

In [54]: 1 pivoted_df.to_csv('data/pivoted_df.csv')

In [37]: 1 pivoted_df.columns

Out[37]: Index(['pitcher', 'season', 'Age', 'Throws', 'Surgery',
'AB_release_speed_weighted_avg', 'CH_release_speed_weighted_avg',
'CS_release_speed_weighted_avg', 'CU_release_speed_weighted_avg',
'EP_release_speed_weighted_avg',
...
'FS_vz0_weighted_avg', 'IN_vz0_weighted_avg', 'KC_vz0_weighted_av
g',
'KN_vz0_weighted_avg', 'PO_vz0_weighted_avg', 'SC_vz0_weighted_av
g',
'SI_vz0_weighted_avg', 'SL_vz0_weighted_avg', 'ST_vz0_weighted_av
g',
'SV_vz0_weighted_avg'],
dtype='object', length=131)

```
In [82]: 1 pd.set_option('display.max_columns', None)
        2 pd.set_option('display.max_rows', None)
```

```
In [91]: 1 pd.reset_option('display.max_columns')
        2 pd.reset_option('display.max_rows')
```

```
In [38]: 1 groovy_df = pivoted_df
```

Drop 'pitcher' column for groovy_df. Only used as ID, should not be necessary.

```
In [40]: 1 groovy_df.drop(columns=['pitcher'], inplace=True)
```


```
In [41]: 1 groovy_df
```

Out[41]:

	season	Age	Throws	Surgery	AB_release_speed_weighted_avg	CH_release_speed_w
--	--------	-----	--------	---------	-------------------------------	--------------------

0	2008	37.0	1	0.0		0.0
1	2009	38.0	1	0.0		0.0
2	2010	39.0	1	0.0		0.0
3	2011	40.0	1	0.0		0.0
4	2012	41.0	1	0.0		0.0
...
3683	2022	25.0	1	0.0		0.0
3684	2023	26.0	1	0.0		0.0
3685	2021	23.0	1	0.0		0.0
3686	2022	24.0	1	0.0		0.0
3687	2023	25.0	1	0.0		0.0

3688 rows × 130 columns



```
In [42]: 1 groovy_df['Surgery'].value_counts()
```

Out[42]:

Surgery	
0.0	2772
1.0	916

Name: count, dtype: int64

```
In [43]: 1 groovy_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3688 entries, 0 to 3687
Columns: 130 entries, season to SV_vz0_weighted_avg
dtypes: float64(128), int64(2)
memory usage: 3.7 MB
```

```
In [101]: ▶ 1 """
2 # Reshape your data as a 2D array of 'pitch_type' column values
3 pitch_type_array = groovy_df['pitch_type'].values.reshape(-1, 1)
4
5 # Fit and transform the 'pitch_type' column to one-hot encoded format
6 pitch_type_ohe = ohe.fit_transform(pitch_type_array)
7
8 # Convert the one-hot encoded result back to a DataFrame
9 pitch_type_df = pd.DataFrame(pitch_type_ohe, columns=ohe.get_feature_names_out(['pitch_type']))
10
11 # Concatenate the new one-hot encoded DataFrame with the original DataFrame (excluding the original 'pitch_type' column)
12 fancy_df = pd.concat([fancy_df.drop('pitch_type', axis=1).reset_index(drop=True), pitch_type_df], axis=1)
13
14 fancy_df.head()
15 """
```

```
Out[101]: "\n# Reshape your data as a 2D array of 'pitch_type' column values\npitch_type_array = groovy_df['pitch_type'].values.reshape(-1, 1)\n\n# Fit and transform the 'pitch_type' column to one-hot encoded format\npitch_type_ohe = ohe.fit_transform(pitch_type_array)\n\n# Convert the one-hot encoded result back to a DataFrame\npitch_type_df = pd.DataFrame(pitch_type_ohe, columns=ohe.get_feature_names_out(['pitch_type']))\n\n# Concatenate the new one-hot encoded DataFrame with the original DataFrame (excluding the original 'pitch_type' column)\nfancy_df = pd.concat([fancy_df.drop('pitch_type', axis=1).reset_index(drop=True), pitch_type_df], axis=1)\n\nfancy_df.head()\n"
```

```
In [44]: ▶ 1 y = groovy_df['Surgery']
2 X = groovy_df.drop('Surgery', axis=1)
```

```
In [45]: ▶ 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```

In [55]: 1 logreg_pipeline = Pipeline([
2         ('scale', StandardScaler()),
3         ('logreg', LogisticRegression(solver='liblinear', max_iter=10000))
4     ])
5
6     # Define the parameter grid to search over
7     param_grid = {
8         'logreg__C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
9         'logreg__penalty': ['l1', 'l2'] # Norm used in the penalization
10    }
11
12    # Initialize GridSearchCV with the pipeline, parameter grid, and desired cv
13    grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='roc_auc')
14
15    # Assuming X_train and y_train are already defined
16    grid_search.fit(X_train, y_train)
17
18    # Best parameters found
19    print("Best parameters: ", grid_search.best_params_)
20
21    # Best cross-validation score
22    print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
23
24    # Test set score using the best parameters
25    print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))

```

Best parameters: {'logreg__C': 10, 'logreg__penalty': 'l1'}

Best cross-validation score: 0.76

Test set score: 0.77

```

In [48]: 1 logreg_pipeline = Pipeline([
2         ('scale', StandardScaler()),
3         ('logreg', LogisticRegression(penalty='l1', C=10.0, solver='liblinear'))
4     ])

```

```

In [49]: 1 logreg_pipeline.fit(X_train, y_train)

```

```

Out[49]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(C=10.0, penalty='l1', solver='liblinear'))])

```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```

In [50]: 1 logreg_pipeline.score(X_test, y_test)

```

```

Out[50]: 0.7687443541102078


```

```

In [51]: 1 y_pred = logreg_pipeline.predict(X_test)

```

```
In [52]: 1 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
2 plt.title('Baseline LR Train Set Classification Metrics')
3 plt.show()
4 print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0.0	0.79	0.93	0.86	827
1.0	0.59	0.28	0.38	280
accuracy			0.77	1107
macro avg	0.69	0.61	0.62	1107
weighted avg	0.74	0.77	0.74	1107

Want more false positives (think needs TJ but doesn't need TJ) than false negatives...

```
In [ ]: 1
```