

```

import numpy as np
import math
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import glob
import os
import io
import re
import unicodedata
from datetime import datetime
from itertools import groupby
from operator import itemgetter
from sklearn.utils.class_weight import compute_class_weight
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, make_scorer
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE
from xgboost import XGBClassifier
from google.colab import files
uploaded = files.upload()

```

No file chosen

pivoted\_df, Logistic Regression Model, baseline

```
pivoted_df = pd.read_csv(io.BytesIO(uploaded['pivoted_df.csv']), index_col=0)
```

```
pivoted_df.head()
```

	season	Age	Throws	Surgery	AB_release_speed_weighted_avg	CH_release_speed_weighted_avg	CS_release_speed_weighted_avg
0	2008	37.0	1	0.0	0.0	82.641530	
1	2009	38.0	1	0.0	0.0	85.012195	
2	2010	39.0	1	0.0	0.0	84.150000	
3	2011	40.0	1	0.0	0.0	83.093750	
4	2012	41.0	1	0.0	0.0	83.001563	

5 rows × 130 columns

```
pivoted_df['Surgery'].value_counts()
```

```

0.0    2772
1.0     916
Name: Surgery, dtype: int64

```

```

y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(solver='liblinear'))
])
```

```
# Define the parameter grid to search over
```

```
param_grid = {
    'logreg__C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
    'logreg__penalty': ['l1', 'l2'] # Norm used in the penalization
}
```

```
# Initialize GridSearchCV with the pipeline, parameter grid, and desired scoring metric
```

```
grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='accuracy')
```

```
# Assuming X_train and y_train are already defined
```

```
grid_search.fit(X_train, y_train)
```

```
# Best parameters found
```

```
print("Best parameters: ", grid_search.best_params_)
```

```
# Best cross-validation score
```

```
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
```


```
# Test set score using the best parameters
```

```
print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
warnings.warn(
Best parameters: {'logreg__C': 10, 'logreg__penalty': 'l1'}
Best cross-validation score: 0.76
Test set score: 0.77
```

```
logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(penalty='l1', C=10.0, solver='liblinear'))
])
```

```
logreg_pipeline.fit(X_train, y_train)
```



```
logreg_pipeline.score(X_test, y_test)
```


```
0.7687443541102078
```

```
y_pred = logreg_pipeline.predict(X_test)
```

```

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('Baseline LR Train Set Classification Metrics')
plt.show()
print(classification_report(y_test, y_pred))

```



	precision	recall	f1-score	support
0.0	0.79	0.93	0.86	827
1.0	0.59	0.28	0.38	280
accuracy			0.77	1107
macro avg	0.69	0.61	0.62	1107
weighted avg	0.74	0.77	0.74	1107

pivoted\_df, Logistic Regression Model, SMOTE

```

smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

```


```
print(pd.Series(y_train_resampled).value_counts())
```

```

1.0    1945
0.0    1945
Name: Surgery, dtype: int64

```

```
logreg_pipeline.fit(X_train_resampled, y_train_resampled)
```

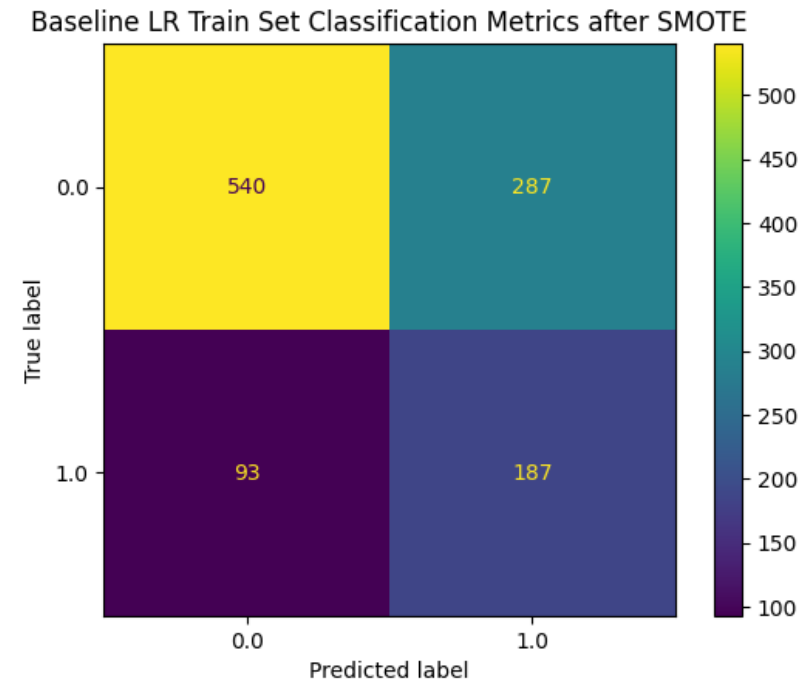


```
y_pred_resampled = logreg_pipeline.predict(X_test)
```

```

ConfusionMatrixDisplay.from_predictions(y_test, y_pred_resampled)
plt.title('Baseline LR Train Set Classification Metrics after SMOTE')
plt.show()
print(classification_report(y_test, y_pred_resampled))

```



	precision	recall	f1-score	support
0.0	0.85	0.65	0.74	827
1.0	0.39	0.67	0.50	280
accuracy			0.66	1107
macro avg	0.62	0.66	0.62	1107
weighted avg	0.74	0.66	0.68	1107

Predicting 1.0 improved, however the rest of the model suffered. Can try condensing the number of features and see how that plays out.

cond\_pivoted\_df, Logistic Regression, Baseline

```
cond_pivoted_df = pd.read_csv(io.BytesIO(uploaded['cond_pivoted_df.csv']), index_col=0)
```

```
cond_pivoted_df.head()
```

	season	Age	Throws	Surgery	BB_release_speed_weighted_avg	FB_release_speed_weighted_avg	OS_release_speed_weighted_avg
0	2008	37.0	1	0.0	75.425843	91.689850	
1	2009	38.0	1	0.0	78.181818	93.479869	
2	2010	39.0	1	0.0	74.666667	93.001617	
3	2011	40.0	1	0.0	76.885714	91.678506	
4	2012	41.0	1	0.0	76.427273	91.965592	

5 rows × 39 columns

```
cond_pivoted_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3688 entries, 0 to 3687
Data columns (total 39 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   season                                     3688 non-null   int64
1   Age                                       3688 non-null   float64
2   Throws                                   3688 non-null   int64
3   Surgery                                   3688 non-null   float64
4   BB_release_speed_weighted_avg           3688 non-null   float64
5   FB_release_speed_weighted_avg           3688 non-null   float64
6   OS_release_speed_weighted_avg           3688 non-null   float64
7   OT_release_speed_weighted_avg           3688 non-null   float64
8   SB_release_speed_weighted_avg           3688 non-null   float64
9   BB_release_pos_x_weighted_avg           3688 non-null   float64
10  FB_release_pos_x_weighted_avg           3688 non-null   float64
11  OS_release_pos_x_weighted_avg           3688 non-null   float64
12  OT_release_pos_x_weighted_avg           3688 non-null   float64
13  SB_release_pos_x_weighted_avg           3688 non-null   float64
14  BB_release_pos_y_weighted_avg           3688 non-null   float64
15  FB_release_pos_y_weighted_avg           3688 non-null   float64
16  OS_release_pos_y_weighted_avg           3688 non-null   float64
17  OT_release_pos_y_weighted_avg           3688 non-null   float64
18  SB_release_pos_y_weighted_avg           3688 non-null   float64
19  BB_release_pos_z_weighted_avg           3688 non-null   float64
20  FB_release_pos_z_weighted_avg           3688 non-null   float64
21  OS_release_pos_z_weighted_avg           3688 non-null   float64
22  OT_release_pos_z_weighted_avg           3688 non-null   float64
23  SB_release_pos_z_weighted_avg           3688 non-null   float64
24  BB_vx0_weighted_avg                     3688 non-null   float64
25  FB_vx0_weighted_avg                     3688 non-null   float64
26  OS_vx0_weighted_avg                     3688 non-null   float64
27  OT_vx0_weighted_avg                     3688 non-null   float64
28  SB_vx0_weighted_avg                     3688 non-null   float64
29  BB_vy0_weighted_avg                     3688 non-null   float64
30  FB_vy0_weighted_avg                     3688 non-null   float64
31  OS_vy0_weighted_avg                     3688 non-null   float64
32  OT_vy0_weighted_avg                     3688 non-null   float64
33  SB_vy0_weighted_avg                     3688 non-null   float64
34  BB_vz0_weighted_avg                     3688 non-null   float64
35  FB_vz0_weighted_avg                     3688 non-null   float64
36  OS_vz0_weighted_avg                     3688 non-null   float64
37  OT_vz0_weighted_avg                     3688 non-null   float64
38  SB_vz0_weighted_avg                     3688 non-null   float64
dtypes: float64(37), int64(2)
memory usage: 1.1 MB
```

```
cond_pivoted_df['Surgery'].value_counts()
```

```
0.0    2772
1.0     916
Name: Surgery, dtype: int64
```

```
cond_pivoted_df.drop(columns=['pitcher'], inplace=True)
```

```
y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```

logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(solver='liblinear'))
])

# Define the parameter grid to search over
param_grid = {
    'logreg__C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization strength
    'logreg__penalty': ['l1', 'l2'] # Norm used in the penalization
}

# Initialize GridSearchCV with the pipeline, parameter grid, and desired scoring metric
grid_search = GridSearchCV(logreg_pipeline, param_grid, cv=5, scoring='accuracy')

# Assuming X_train and y_train are already defined
grid_search.fit(X_train_resampled, y_train_resampled)

# Best parameters found
print("Best parameters: ", grid_search.best_params_)

# Best cross-validation score
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))

# Test set score using the best parameters
print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))

/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
  warnings.warn(
Best parameters: {'logreg__C': 10, 'logreg__penalty': 'l1'}
Best cross-validation score: 0.67
Test set score: 0.64


```

```

logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(penalty='l1', C=10.0, solver='liblinear'))
])

```

```
logreg_pipeline.fit(X_train, y_train)
```



```
logreg_pipeline.score(X_test, y_test)
```

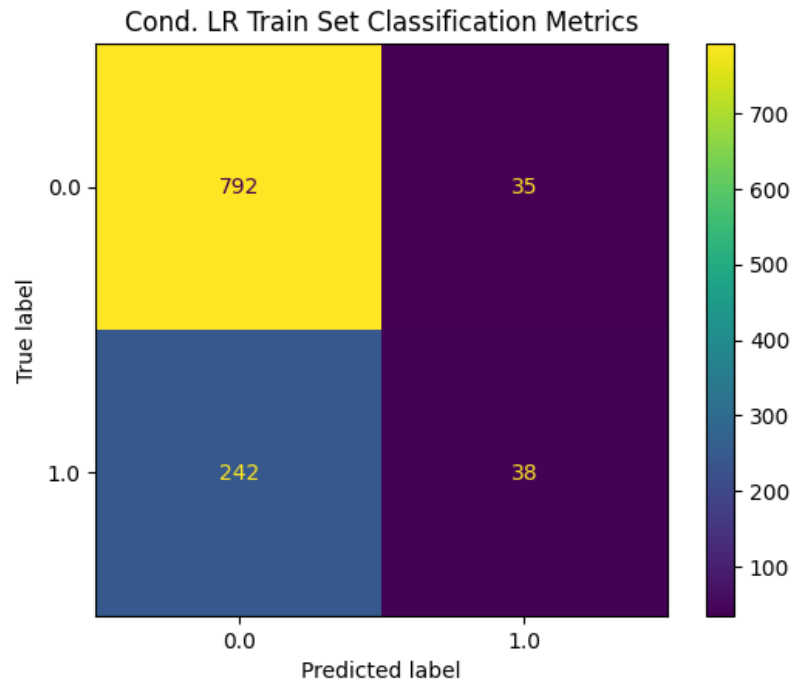
```
0.7497741644083108
```

```
y_pred = logreg_pipeline.predict(X_test)
```

```

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('Cond. LR Train Set Classification Metrics')
plt.show()
print(classification_report(y_test, y_pred))

```



	precision	recall	f1-score	support
0.0	0.77	0.96	0.85	827
1.0	0.52	0.14	0.22	280
accuracy			0.75	1107
macro avg	0.64	0.55	0.53	1107
weighted avg	0.70	0.75	0.69	1107

Score for 1.0 predicting Surgery too low. Now try with SMOTE.


cond\_pivoted\_df, Logistic Regression Model, SMOTE

```
smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

print(pd.Series(y_train_resampled).value_counts())
```

```
1.0    1945
0.0    1945
Name: Surgery, dtype: int64
```


```
logreg_pipeline.fit(X_train_resampled, y_train_resampled)
```



```
y_pred_resampled = logreg_pipeline.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_resampled)
plt.title('Cond. Baseline LR Train Set Classification Metrics after SMOTE')
plt.show()
print(classification_report(y_test, y_pred_resampled))
```

Cond. Baseline LR Train Set Classification Metrics after SMOTE



	precision	recall	f1-score	support
0.0	0.83	0.64	0.73	827
1.0	0.37	0.62	0.47	280
accuracy			0.64	1107
macro avg	0.60	0.63	0.60	1107
weighted avg	0.72	0.64	0.66	1107

pivot\_df is better at predicting TP & TN. cond\_pivot\_df is better at FP & FN. Trade-off between two models is slim. SMOTE helped, can I address class imbalance further, or will that lead to issues? Still need to see improvement. Need to try PCA, Random Forest, XG Boost

Try to see how GridSearch works when I add in class balance. Will try this out for both pivoted\_df and cond\_pivoted\_df

pivoted\_df, Logistic Regression Model, adjust class weights

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```



```
# Define the parameter grid to search over, including class weights
class_weights = [None, 'balanced', {0: 1, 1: 2}, {0: 1, 1: 3}, {0: 1, 1: 4}, {0: 1, 1: 5}]
param_grid = {
    'logreg__C': [0.01, 0.1, 1, 10],
    'logreg__penalty': ['l1', 'l2'],
    'logreg__class_weight': class_weights,
    'logreg__max_iter': [5000],
    'logreg__tol': [0.01]
}

# Create a scoring function that focuses on recall for the positive class
recall_scorer = make_scorer(recall_score, pos_label=1)

# Initialize GridSearchCV with the pipeline, parameter grid, and recall as the scoring metric
grid_search = GridSearchCV(weight_logreg_pipeline, param_grid, cv=5, scoring=recall_scorer, n_jobs=-1)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)


# Print the best parameters found and the best recall score
print("Best parameters: ", grid_search.best_params_)
print("Best cross-validation recall score: {:.2f}".format(grid_search.best_score_))

# Evaluate the best model on the test set
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
print("Test set recall score: {:.2f}".format(recall_score(y_test, y_pred)))

Best parameters: {'logreg__C': 0.01, 'logreg__class_weight': {0: 1, 1: 5}, 'logreg__max_iter': 5000, 'logreg__penal
Best cross-validation recall score: 0.87
Test set recall score: 0.85
```

```
weight_logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(penalty='l1', C=0.01, class_weight={0:1, 1:5}, solver='liblinear'))
])
```

```
weight_logreg_pipeline.fit(X_train, y_train)
```




```
weight_logreg_pipeline.score(X_test, y_test)
```

```
0.5167118337850045
```

```
y_pred = weight_logreg_pipeline.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('Weighted LR Train Set Classification Metrics')
plt.show()
print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0.0	0.89	0.40	0.55	827
1.0	0.33	0.85	0.47	280
accuracy			0.52	1107
macro avg	0.61	0.63	0.51	1107
weighted avg	0.75	0.52	0.53	1107

This looks better. The number of False negatives has decreased dramatically. Would like to see further improvement to be more accurate all around. Now to test `cond_pivoted_df`

`cond_pivoted_df`, Logistic Regression Model, adjust class weights.

```
y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Define the parameter grid to search over, including class weights
class_weights = [None, 'balanced', {0: 1, 1: 2}, {0: 1, 1: 3}, {0: 1, 1: 4}]
param_grid = {
    'logreg__C': [0.01, 0.1, 1, 10, 100],
    'logreg__penalty': ['l1', 'l2'],
    'logreg__class_weight': class_weights,
    'logreg__max_iter': [5000],
    'logreg__tol': [0.01]
}

# Create a scoring function that focuses on recall for the positive class
recall_scorer = make_scorer(recall_score, pos_label=1)

# Initialize GridSearchCV with the pipeline, parameter grid, and recall as the scoring metric
grid_search = GridSearchCV(weight_logreg_pipeline, param_grid, cv=5, scoring=recall_scorer)

# Fit the grid search to the data
grid_search.fit(X_train_resampled, y_train_resampled)

# Print the best parameters found and the best recall score
print("Best parameters: ", grid_search.best_params_)
print("Best cross-validation recall score: {:.2f}".format(grid_search.best_score_))

# Evaluate the best model on the test set
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
print("Test set recall score: {:.2f}".format(recall_score(y_test, y_pred)))

Best parameters: {'logreg__C': 0.01, 'logreg__class_weight': {0: 1, 1: 4}, 'logreg__max_iter': 5000, 'logreg__penal
Best cross-validation recall score: 1.00
Test set recall score: 0.99
```


```
smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

print(pd.Series(y_train_resampled).value_counts())

1.0    1945
0.0    1945
Name: Surgery, dtype: int64

weight_logreg_pipeline = Pipeline([
    ('scale', StandardScaler()),
    ('logreg', LogisticRegression(penalty='l1', C=0.01, class_weight={0:1, 1:3}, solver='liblinear'))
])
```

```
weight_logreg_pipeline.fit(X_train_resampled, y_train_resampled)
```



```
y_pred_resampled = weight_logreg_pipeline.predict(X_test)

ConfusionMatrixDisplay.from_predictions(y_test, y_pred_resampled)
plt.title('Weight Cond. LR Train Set Classification Metrics after SMOTE')
plt.show()
print(classification_report(y_test, y_pred_resampled))
```




```

0.          , 0.          , 0.          , 0.          , 0.          ,
0.02142519, 0.          , 0.          , 0.          , 0.06182032,
0.          , 0.04426144, 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.04978206, 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.06168103, 0.          , 0.          , 0.          , 0.          ,
0.          , 0.0381441 , 0.          , 0.          , 0.          ,
0.0121275 , 0.          , 0.02063147, 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.04179382,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
]
```

```


def plot_feature_importances(model, n_top_features=20):
    importances = model.feature_importances_
    indices = np.argsort(importances)[-n_top_features:]
    plt.figure(figsize=(10,10))
    plt.title('Top N Feature Importances')
    plt.barh(range(n_top_features), importances[indices], align='center')
    plt.yticks(range(n_top_features), [X_train.columns[i] for i in indices])
    plt.xlabel('Feature importance')
    plt.ylabel('Feature')
    plt.ylim(-1, n_top_features)

plot_feature_importances(tree_clf, n_top_features=20)
plt.show()
```



```
pred = tree_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)  
plt.title('pivoted_df Decision Tree Classifier')  
plt.show()  
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.77	0.98	0.86	827
1.0	0.65	0.11	0.19	280
accuracy			0.76	1107
macro avg	0.71	0.55	0.53	1107
weighted avg	0.74	0.76	0.69	1107

```

param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

tree_clf = DecisionTreeClassifier()
scorer = make_scorer(recall_score)
grid_search = GridSearchCV(estimator=tree_clf, param_grid=param_grid, scoring=scorer, cv=5, n_jobs=-1)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Print the best parameters and the best score
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

# Evaluate the best model found by GridSearchCV on the test set
best_tree = grid_search.best_estimator_
y_pred = best_tree.predict(X_test)
print("Test set accuracy:", accuracy_score(y_test, y_pred))

Best parameters: {'class_weight': 'balanced', 'criterion': 'gini', 'max_depth': 5, 'min_samples_leaf': 4, 'min_sampl
Best score: 0.6540600393700788
Test set accuracy: 0.5817524841915086

```

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
tree_clf = DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=1, min_samples_split=10)
tree_clf.fit(X_train, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_split=10)
```

```
pred = tree_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('Grid Search Param Decision Tree Classifier')
plt.show()
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.77	0.94	0.85	827
1.0	0.53	0.19	0.27	280
accuracy			0.75	1107
macro avg	0.65	0.56	0.56	1107
weighted avg	0.71	0.75	0.70	1107

Grid Search improved on TN & FN but worse on TP & TN. Would rather improve on TN than FN. Can try rerunning.

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```


```
tree_clf = DecisionTreeClassifier(criterion='gini', class_weight='balanced', max_depth=5, min_samples_leaf=4, min_sample
tree_clf.fit(X_train, y_train)
```



```
DecisionTreeClassifier
DecisionTreeClassifier(class_weight='balanced', max_depth=5, min_samples_leaf=4,
min_samples_split=10)
```

```
pred = tree_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('Grid Search Recall Decision Tree Classifier')
plt.show()
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.87	0.52	0.65	827
1.0	0.35	0.76	0.48	280
accuracy			0.58	1107
macro avg	0.61	0.64	0.56	1107
weighted avg	0.74	0.58	0.61	1107

cond\_pivoted\_df, Decision Tree Classifier

```
cond_pivoted_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3688 entries, 0 to 3687
Data columns (total 39 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   season                                     3688 non-null   int64
1   Age                                       3688 non-null   float64
2   Throws                                   3688 non-null   int64
3   Surgery                                   3688 non-null   float64
4   BB_release_speed_weighted_avg            3688 non-null   float64
5   FB_release_speed_weighted_avg            3688 non-null   float64
6   OS_release_speed_weighted_avg            3688 non-null   float64
7   OT_release_speed_weighted_avg            3688 non-null   float64
8   SB_release_speed_weighted_avg            3688 non-null   float64
9   BB_release_pos_x_weighted_avg            3688 non-null   float64
10  FB_release_pos_x_weighted_avg            3688 non-null   float64
```

```

11 OS_release_pos_x_weighted_avg 3688 non-null float64
12 OT_release_pos_x_weighted_avg 3688 non-null float64
13 SB_release_pos_x_weighted_avg 3688 non-null float64
14 BB_release_pos_y_weighted_avg 3688 non-null float64
15 FB_release_pos_y_weighted_avg 3688 non-null float64
16 OS_release_pos_y_weighted_avg 3688 non-null float64
17 OT_release_pos_y_weighted_avg 3688 non-null float64
18 SB_release_pos_y_weighted_avg 3688 non-null float64
19 BB_release_pos_z_weighted_avg 3688 non-null float64
20 FB_release_pos_z_weighted_avg 3688 non-null float64
21 OS_release_pos_z_weighted_avg 3688 non-null float64
22 OT_release_pos_z_weighted_avg 3688 non-null float64
23 SB_release_pos_z_weighted_avg 3688 non-null float64
24 BB_vx0_weighted_avg            3688 non-null float64
25 FB_vx0_weighted_avg            3688 non-null float64
26 OS_vx0_weighted_avg            3688 non-null float64
27 OT_vx0_weighted_avg            3688 non-null float64
28 SB_vx0_weighted_avg            3688 non-null float64
29 BB_vy0_weighted_avg            3688 non-null float64
30 FB_vy0_weighted_avg            3688 non-null float64
31 OS_vy0_weighted_avg            3688 non-null float64
32 OT_vy0_weighted_avg            3688 non-null float64
33 SB_vy0_weighted_avg            3688 non-null float64
34 BB_vz0_weighted_avg            3688 non-null float64
35 FB_vz0_weighted_avg            3688 non-null float64
36 OS_vz0_weighted_avg            3688 non-null float64
37 OT_vz0_weighted_avg            3688 non-null float64
38 SB_vz0_weighted_avg            3688 non-null float64
dtypes: float64(37), int64(2)
memory usage: 1.1 MB

```

```

y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```

tree_clf = DecisionTreeClassifier(criterion='gini', max_depth=5)
tree_clf.fit(X_train, y_train)

```

```


▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)

```

```

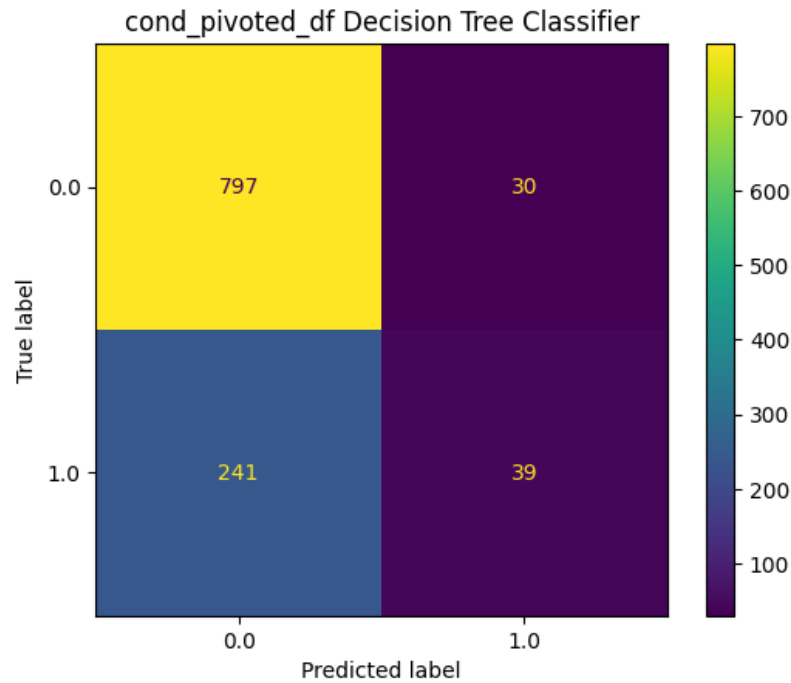
plot_feature_importances(tree_clf, n_top_features=20)
plt.show()

```



```
pred = tree_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('cond_pivoted_df Decision Tree Classifier')
plt.show()
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.77	0.96	0.85	827
1.0	0.57	0.14	0.22	280
accuracy			0.76	1107
macro avg	0.67	0.55	0.54	1107
weighted avg	0.72	0.76	0.70	1107

cond\_pivoted\_df not predicting TP & FP well.

```

param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

tree_clf = DecisionTreeClassifier()
scorer = make_scorer(recall_score)
grid_search = GridSearchCV(estimator=tree_clf, param_grid=param_grid, scoring=scorer, cv=5, n_jobs=-1)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Print the best parameters and the best score
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

# Evaluate the best model found by GridSearchCV on the test set
best_tree = grid_search.best_estimator_
y_pred = best_tree.predict(X_test)
print("Test set accuracy:", accuracy_score(y_test, y_pred))

```

```

Best parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': 5, 'min_samples_leaf': 4, 'min_sa
Best score: 0.7031373031496063
Test set accuracy: 0.6621499548328816

```

```
y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)


tree_clf = DecisionTreeClassifier(class_weight='balanced', criterion='entropy', max_depth=5, min_samples_leaf=4, min_sam
tree_clf.fit(X_train, y_train)
```

▼ DecisionTreeClassifier

DecisionTreeClassifier(class\_weight='balanced', criterion='entropy', max\_depth=5, min\_samples\_leaf=4, min\_samples\_split=10)

```
pred = tree_clf.predict(X_test)

ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('cond_pivoted_df recall Grid Search Param DecTreeClass')
plt.show()
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.83	0.69	0.75	827
1.0	0.39	0.59	0.47	280
accuracy			0.66	1107
macro avg	0.61	0.64	0.61	1107
weighted avg	0.72	0.66	0.68	1107

Better, but FN is high compared to other models.

Start coding or [generate](#) with AI.

pivoted\_df, Random Forest

```
y = pivoted_df['Surgery']
```

```
x = pivoted_ar.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
forest_clf = RandomForestClassifier(n_estimators=100, max_depth=5)
forest_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(max_depth=5)
```


```
forest_clf.score(X_train, y_train)
```

```
0.7667570709027509
```

```
forest_clf.score(X_test, y_test)
```

```
0.7515808491418248
```

```
plot_feature_importances(forest_clf, n_top_features=20)
plt.show()
```




```
forest2_clf = RandomForestClassifier(n_estimators=5, max_features=10, max_depth=2)
forest2_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(max_depth=2, max_features=10, n_estimators=5)
```


```
rf_tree_1 = forest2_clf.estimators_[0]
```

```
plot_feature_importances(rf_tree_1, n_top_features=20)
plt.show()
```



```
rf_tree_2 = forest2_clf.estimators_[1]
```

```
plot_feature_importances(rf_tree_2, n_top_features=20)
plt.show()
```



```
param_grid_rf = {
    'n_estimators': [100, 200, 300],
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

forest_clf = RandomForestClassifier()
accuracy_scorer = make_scorer(accuracy_score)
grid_search_rf = GridSearchCV(estimator=forest_clf, param_grid=param_grid_rf, scoring=accuracy_scorer, cv=5, n_jobs=-1)
grid_search_rf.fit(X_train, y_train)

print("Best parameters:", grid_search_rf.best_params_)
print("Best accuracy score:", grid_search_rf.best_score_)

best_rf = grid_search_rf.best_estimator_
y_pred_rf = best_rf.predict(X_test)

print("Test set accuracy score:", accuracy_score(y_test, y_pred_rf))
```

```
Best parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 4, 'min
Best accuracy score: 0.8109314320843268
```



Test set accuracy score: 0.8175248419150858

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
forest_clf = RandomForestClassifier(class_weight='balanced', criterion='entropy', max_depth=None, min_samples_leaf=4, min_samples_split=10)
forest_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(class_weight='balanced', criterion='entropy',
                        min_samples_leaf=4, n_estimators=300)
```

```
pred = forest_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('pivoted_df accuracy Grid Search Param ForestClf')
plt.show()
print(classification_report(y_test, pred))
```



Again, TN & TP are good, but FP is too small compared to FN. Will try this with recall score.

```

param_grid_rf = {
    'n_estimators': [100, 200, 300],
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

forest_clf = RandomForestClassifier()
recall_scorer = make_scorer(recall_score)
grid_search_rf = GridSearchCV(estimator=forest_clf, param_grid=param_grid_rf, scoring=accuracy_scorer, cv=5, n_jobs=-1)
grid_search_rf.fit(X_train, y_train)

print("Best parameters:", grid_search_rf.best_params_)
print("Best accuracy score:", grid_search_rf.best_score_)

best_rf = grid_search_rf.best_estimator_
y_pred_rf = best_rf.predict(X_test)

print("Test set accuracy score:", recall_score(y_test, y_pred_rf))

Best parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 2, 'min
Best accuracy score: 0.810923935045657
Test set accuracy score: 0.34285714285714286

```

```

y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

```

forest_clf = RandomForestClassifier(class_weight='balanced', criterion='entropy', max_depth=None, min_samples_leaf=2, mi
forest_clf.fit(X_train, y_train)

```

```

▼ RandomForestClassifier
RandomForestClassifier(class_weight='balanced', criterion='entropy',
                        min_samples_leaf=2, min_samples_split=10)


```

```

pred = forest_clf.predict(X_test)

ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('pivoted_df recall Grid Search Param ForestClf')
plt.show()
print(classification_report(y_test, pred))

```



	precision	recall	f1-score	support
0.0	0.82	0.97	0.89	827
1.0	0.82	0.36	0.50	280
accuracy			0.82	1107
macro avg	0.82	0.67	0.69	1107
weighted avg	0.82	0.82	0.79	1107

Not great. Would like to see lower FN and higher TP & FP.

cond\_pivoted\_df, Random Forest

```
y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
forest_clf = RandomForestClassifier(n_estimators=100, max_depth=5)
forest_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(max_depth=5)
```


```
forest_clf.score(X_train, y_train)
```

```
0.7729562185199536
```

```
forest_clf.score(X_test, y_test)
```

```
0.7497741644083108
```

```
plot_feature_importances(forest_clf, n_top_features=20)
plt.show()
```



```

param_grid_rf = {
    'n_estimators': [100, 200, 300],
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

forest_clf = RandomForestClassifier()
accuracy_scorer = make_scorer(accuracy_score)
grid_search_rf = GridSearchCV(estimator=forest_clf, param_grid=param_grid_rf, scoring=accuracy_scorer, cv=5, n_jobs=-1)
grid_search_rf.fit(X_train, y_train)

print("Best parameters:", grid_search_rf.best_params_)
print("Best accuracy score:", grid_search_rf.best_score_)

best_rf = grid_search_rf.best_estimator_
y_pred_rf = best_rf.predict(X_test)

print("Test set accuracy score:", accuracy_score(y_test, y_pred_rf))

t parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': 15, 'min_samples_leaf': 2, 'min_sampl
t accuracy score: 0.8085953548348404

```

t set accuracy score: 0.8121047877145439

```
y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)
```


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
forest_clf = RandomForestClassifier(class_weight='balanced', criterion='entropy', max_depth=15, min_samples_leaf=2, min_
forest_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(class_weight='balanced', criterion='entropy',
                        max_depth=15, min_samples_leaf=2, min_samples_split=5,
                        n_estimators=200)
```

```
pred = forest_clf.predict(X_test)
```

```
ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('cond_pivoted_df accuracy Grid Search Param ForestClf')
plt.show()
print(classification_report(y_test, pred))
```



	precision	recall	f1-score	support
0.0	0.81	0.98	0.89	827
1.0	0.83	0.33	0.47	280
accuracy			0.81	1107
macro avg	0.82	0.65	0.68	1107
weighted avg	0.81	0.81	0.78	1107

Try recall score.

```

param_grid_rf = {
    'n_estimators': [100, 200, 300],
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'class_weight': [None, 'balanced', {0:1, 1:2}, {0:1, 1:3}]
}

forest_clf = RandomForestClassifier()
recall_scorer = make_scorer(recall_score)
grid_search_rf = GridSearchCV(estimator=forest_clf, param_grid=param_grid_rf, scoring=accuracy_scorer, cv=5, n_jobs=-1)
grid_search_rf.fit(X_train, y_train)

print("Best parameters:", grid_search_rf.best_params_)
print("Best accuracy score:", grid_search_rf.best_score_)

best_rf = grid_search_rf.best_estimator_
y_pred_rf = best_rf.predict(X_test)

print("Test set accuracy score:", recall_score(y_test, y_pred_rf))

t parameters: {'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': 15, 'min_samples_leaf': 2, 'min_sampl
t accuracy score: 0.8074348132487668
t set accuracy score: 0.33214285714285713

```

```

y = cond_pivoted_df['Surgery']
X = cond_pivoted_df.drop('Surgery', axis=1)

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

```

forest_clf = RandomForestClassifier(class_weight='balanced', criterion='entropy', max_depth=15, min_samples_leaf=2, min_
forest_clf.fit(X_train, y_train)

```

```

▼                                RandomForestClassifier
RandomForestClassifier(class_weight='balanced', criterion='entropy',
                        max_depth=15, min_samples_leaf=2, min_samples_split=5)


```

```

pred = forest_clf.predict(X_test)

ConfusionMatrixDisplay.from_predictions(y_test, pred)
plt.title('cond_pivoted_df recall Grid Search Param ForestClf')
plt.show()
print(classification_report(y_test, pred))

```



	precision	recall	f1-score	support
0.0	0.81	0.97	0.88	827
1.0	0.80	0.33	0.46	280
accuracy			0.81	1107
macro avg	0.80	0.65	0.67	1107
weighted avg	0.81	0.81	0.78	1107

Odd. Scores are very similar. Is there a way to better these scores? May just have to move onto another model. XG Boost?

pivoted\_df XG Boost

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
xgb_clf = XGBClassifier(use_label_encoder=False, eval_metric='logloss')
```

```
xgb_clf.fit(X_train, y_train)
```

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, gamma=None, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_bin=None, max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None, n_estimators=None,
               n_jobs=None, num_parallel_tree=None, random_state=None, ...)

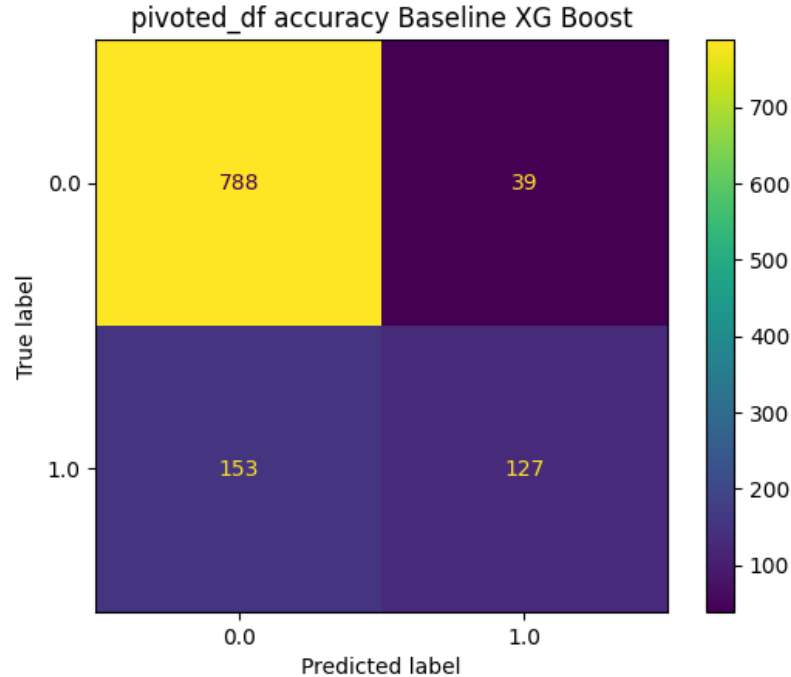
```

```
y_pred = xgb_clf.predict(X_test)
```

```

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('pivoted_df accuracy Baseline XG Boost')
plt.show()
print(classification_report(y_test, pred))

```



	precision	recall	f1-score	support
0.0	0.81	0.97	0.88	827
1.0	0.80	0.33	0.46	280
accuracy			0.81	1107
macro avg	0.80	0.65	0.67	1107
weighted avg	0.81	0.81	0.78	1107

Pretty good. Would like to see lower FN and higher scores for everything else. Will try GridSearch.

```

param_grid_xgb = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1],
    'colsample_bytree': [0.8, 1]
}

xgb_clf = XGBClassifier(objective='binary:logistic', use_label_encoder=False, eval_metric='logloss')

grid_search_xgb = GridSearchCV(estimator=xgb_clf, param_grid=param_grid_xgb, scoring='recall', cv=5, n_jobs=-1)
grid_search_xgb.fit(X_train, y_train)

print("Best parameters:", grid_search_xgb.best_params_)
print("Best recall score:", grid_search_xgb.best_score_)

best_xgb = grid_search_xgb.best_estimator_
y_pred_best_xgb = best_xgb.predict(X_test)
print("Test set recall score:", recall_score(y_test, y_pred_best_xgb))
print(classification_report(y_test, y_pred_best_xgb))

Best parameters: {'colsample_bytree': 1, 'learning_rate': 0.2, 'max_depth': 7, 'n_estimators': 200, 'subsample': 0.8}
Best recall score: 0.4323941929133858
Test set recall score: 0.45357142857142857

```



	precision	recall	f1-score	support
0.0	0.84	0.96	0.89	827
1.0	0.79	0.45	0.58	280
accuracy			0.83	1107
macro avg	0.81	0.71	0.74	1107
weighted avg	0.83	0.83	0.81	1107

```
y = pivoted_df['Surgery']
X = pivoted_df.drop('Surgery', axis=1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
xgb_clf = XGBClassifier(colsample_bytree=1, learning_rate=0.2, max_depth=7, n_estimators=200, subsample=0.8, use_label_e
```

```
xgb_clf.fit(X_train, y_train)
```

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None, colsample_bytree=1,
               device=None, early_stopping_rounds=None, enable_categorical=False,
               eval_metric='logloss', feature_types=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.2, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=7, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=200, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...)

```

```
y_pred = xgb_clf.predict(X_test)
```

```

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('pivoted_df recall GridSearch XG Boost')
plt.show()
print(classification_report(y_test, pred))

```

